

Huanhuan Ma

CONTACT INFORMATION	 huanhuanma.top  hma42@uic.edu  +1 773 827 4356	 Google Scholar	 GitHub  in LinkedIn  Chicago, USA
EDUCATION	University of Illinois Chicago , Chicago, USA PhD in Computer Science, expected May 2029 - Advisor: Prof. Philip S. Yu (ACM/IEEE/AAAS Fellow) University of Chinese Academy of Sciences , Beijing, China Master in Artificial Intelligence - Advisors: Prof. Liang Wang, Prof. Shu Wu, Prof. Qiang Liu - Thesis: Explainable Automated Fact Verification Zhengzhou University , Zhengzhou, China Bachelor in Software Engineering	Aug 2025 – Present Sep 2021 – Jul 2024 Sep 2016 – Jul 2020	
RESEARCH CONTRIBUTIONS	My research develops evaluation frameworks and personalization methods for large language models (LLMs) so that AI systems serve individual users in ways that are trustworthy, fair, and reliable. Published work has received over 236 citations (Google Scholar). Trustworthy LLMs & Fact Verification: built EX-FEVER [2], a multi-hop explainable fact-verification benchmark, and LOGRAN [3], a soft-logic regularization method for multimodal out-of-context detection. Reliable Behavioral Evaluation: proposed the Core Sentiment Inventory (CSI) [7], an implicit behavioral probe that bypasses self-report bias to test whether LLMs behave consistently. Personalization & User Alignment: developed an evaluation framework showing that anti-sycophancy methods can impair rational belief updating [1].		
FIRST-AUTHOR PUBLICATIONS	* = equal contribution. [1] Huanhuan Ma , Henry Peng Zou, Chengze Li, Enze Ma, Yunyue Su, Philip S. Yu. “Sycophancy Suppression Can Impair Rational Updating: Anti-Sycophancy Should Preserve the Ability to Update.” <i>Findings of the Association for Computational Linguistics (EMNLP 2026 Findings)</i> . [OpenReview] [2] Huanhuan Ma , Weizhi Xu, Yifan Wei, et al. “EX-FEVER: A Dataset for Multi-hop Explainable Fact Verification.” <i>Findings of the Association for Computational Linguistics (ACL 2024 Findings)</i> . [PDF], [Code/Dataset] [3] Huanhuan Ma [*] , Jinghao Zhang [*] , Qiang Liu, Shu Wu, Liang Wang. “Interpretable Multi-modal Out-of-Context Detection with Soft Logic Regularization.” <i>Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP 2024, Oral)</i> . [PDF]		
OTHER PUBLICATIONS	[4] Yifan Wei, Xiaoyan Yu, Yixuan Weng, Huanhuan Ma , et al. “Does Knowledge Localization Hold True? Surprising Differences Between Entity and Relation Perspectives in Language Models.” <i>Proceedings of the 33rd ACM International Conference on Information and Knowledge Management (CIKM 2024)</i> . [PDF] [5] Yifan Wei, Yisong Su, Huanhuan Ma , et al. “MenatQA: A New Dataset for Testing the Temporal Comprehension and Reasoning Abilities of Large Language Models.” <i>Findings of the Association for Computational Linguistics (EMNLP 2023 Findings)</i> . [PDF] [6] Liping Wang, Qiang Liu, Huanhuan Ma , Shu Wu, Liang Wang. “Multi-Cause Learning for Diagnosis Prediction.” <i>Proceedings of the International Conference on Data Mining and Big Data (DMBD 2022)</i> . [PDF]		

PREPRINTS &
UNDER REVIEW

The following are preprints or submitted manuscripts, not yet peer-reviewed.

- [7] **Huanhuan Ma**, Haisong Gong, Xiaoyuan Yi, Xing Xie, Philip S. Yu, Dongkuan Xu. "Beyond BFI: The CSI for Enhanced Reliability and Validity in Evaluating LLM Personality Traits." Preprint. [arXiv]
- [8] Hanrong Zhang, Yankai Chen, Shicheng Fan, Dehai Min, Shaowen Chen, **Huanhuan Ma**, et al. "Scaling LLM Agent Learning with Data Synthesis: A Comprehensive Survey." Under review at EMNLP 2026.
- [9] Enze Ma, Yufan Zhou, Jie Yang, **Huanhuan Ma**, et al. "MEMPROBE: Probing Long-Term Agent Memory via Hidden User-State Recovery." Under review at NeurIPS 2026.
- [10] Weizhi Zhang, Wei-Chieh Huang, Yueqing Liang, et al., **Huanhuan Ma**, et al., Kai Shu, Xue Liu, Philip S. Yu. "JARVIS-Bench: Benchmarking Personal Intelligence Agents on Long-Horizon Real-User Daily Traces." Under review at NeurIPS 2026.
- [11] Henry Peng Zou, Chunyu Miao, Wei-Chieh Huang, et al., **Huanhuan Ma**, et al., Xue Liu, Philip S. Yu. "InterruptBench: Evaluating LLM Agents Under User Interruptions in Long-Horizon Web Tasks." Preprint. [arXiv]
- [12] Haisong Gong, **Huanhuan Ma**, Qiang Liu, Shu Wu, Liang Wang. "Navigating the Noisy Crowd: Finding Key Information for Claim Verification." Preprint. [arXiv]
- [13] Shuhao Gu, Jialing Zhang, et al., **Huanhuan Ma**, et al., Xinlong Wang, Guang Liu. "Infinity-MM: Scaling Multimodal Performance with Large-Scale and High-Quality Instruction Data." Preprint. [arXiv]

RESEARCH
EXPERIENCE

Personalization & User Alignment

University of Illinois Chicago, Chicago, USA

Research Assistant, Advisor: *Prof. Philip S. Yu*

Aug 2025 – Present

- Studied how LLM applications can personalize for different users while preserving rational belief updating and avoiding sycophancy from over-personalization [1]. Findings of EMNLP 2026.

Behavioral Evaluation for LLMs

North Carolina State University (Remote)

Collaborator: *Dr. Dongkuan Xu*

Dec 2023 – Dec 2024

- Proposed the Core Sentiment Inventory (CSI) [7] for reliable behavioral evaluation of LLMs.

Trustworthy LLMs & Fact Verification

University of Chinese Academy of Sciences

Master's Student, Advisors: *Prof. Liang Wang, Prof. Shu Wu, Prof. Qiang Liu*

Aug 2021 – Jul 2024

- Created **EX-FEVER** (ACL 2024) [2], a benchmark for multi-hop explainable fact verification.
- Developed **LOGRAN** for multimodal out-of-context detection (ICASSP 2024 Oral) [3].
- Contributed to **EACON** [12], an evidence-abstraction and claim-deconstruction framework for LLM-based claim verification over noisy evidence.

INTERNSHIP

Research Intern

Beijing Academy of Artificial Intelligence (BAAI)

Mentor: *Dr. Li Du*

Jul 2024 – Apr 2025

- Engineered instruction-data pipelines (collection, quality scoring, deduplication, difficulty filtering) for **Infinity-MM** [13], a **40M-sample** open multimodal instruction dataset.
- Designed tag-conditioned **synthetic data generation** targeting diagnosed failure modes; the resulting corpus trained **Aquila-VL-2B** to state-of-the-art performance at the 2B scale.

AWARDS

- **NSF ACCESS Explore Award (Principal Investigator)**, 200,000 ACCESS computing credits 2026–2027
- **ACL ARR Outstanding Reviewer Award** (July 2025 Cycle) 2025
- Ministry of Education – Huawei "Intelligent Base" Future Star 2022

ACADEMIC
SERVICE

Reviewer:

- ICLR 2025, 2026
- NeurIPS 2026
- ACL ARR 2025, 2026
- CIKM 2024, 2025
- AAAI 2026 Workshop (PerFM)

OPEN SOURCE

Awesome-LLM-based-Evaluators

[\[GitHub\]](#) ★33

Curated collection of state-of-the-art research on **LLM-as-a-Judge** and behavioral-evaluation.